

Spoken Digit Recognition

In this notebook, we will work on Spoken Digit Recognition.

Input - speech signal, output - digit number

1. Reading the dataset. and Preprocess the data set.
2. Training the LSTM with RAW data
3. Converting to spectrogram and Training the LSTM network
4. Creating the augmented data and doing step 2 and 3 again.

```
In [1]: 1 from google.colab import drive
        2 drive.mount('/gdrive')
        3 %cd /gdrive
```

Mounted at /gdrive
/gdrive

```
In [2]: 1 import numpy as np
        2 import pandas as pd
        3 import librosa
        4 import os
        5 from tqdm import tqdm
        6 from sklearn.model_selection import train_test_split
        7 import matplotlib.pyplot as plt
        8 import tensorflow as tf
        9
       10 ##if you need any imports you can do that here.
```

We shared recordings.zip, please unzip those.

```
In [3]: 1 #read the all file names in the recordings folder given by us
2 #(if you get entire path, it is very useful in future)
3 #save those files names as list in "all_files"
4
5 #!unzip  "/gdrive/My Drive/Spoken_digit/recordings.zip" -d "/gdrive/My
6
7 all_files_name = os.listdir('/gdrive/My Drive/recordings')
8
9 all_files = []
10 labels = []
11 # get path of all_files
12 for i in tqdm(all_files_name):
13     file_path = '/gdrive/My Drive/recordings/'+str(i)
14     all_files.append(file_path)
15     split = int(i.split('_')[0])
16     labels.append(split)
17
```

100%|██████████| 2000/2000 [00:00<00:00, 436611.04it/s]

```
In [4]: 1 all_files[:5]
```

```
Out[4]: ['/gdrive/My Drive/recordings/7_theo_33.wav',
'/gdrive/My Drive/recordings/5_jackson_35.wav',
'/gdrive/My Drive/recordings/3_jackson_19.wav',
'/gdrive/My Drive/recordings/9_yweweler_20.wav',
'/gdrive/My Drive/recordings/0_theo_28.wav']
```

```
In [5]: 1 labels[:5]
```

```
Out[5]: [7, 5, 3, 9, 0]
```

Grader function 1

```
In [6]: 1 def grader_files():
2     temp = len(all_files)==2000
3     temp1 = all([x[-3:]=="wav" for x in all_files])
4     temp = temp and temp1
5     return temp
6 grader_files()
```

```
Out[6]: True
```

Create a dataframe(name=df_audio) with two columns(path, label).

You can get the label from the first letter of name.

Eg: 0_jackson_0 --> 0

0_jackson_43 --> 0

```
In [7]: 1 #Create a dataframe(name=df_audio) with two columns(path, label).
2 #You can get the label from the first letter of name.
3 #Eg: 0_jackson_0 --> 0
4 #0_jackson_43 --> 0
5 df_audio = pd.DataFrame({'path' :all_files, 'label':labels })
6
```

In [8]:

```
1 #info
2 df_audio.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 2000 entries, 0 to 1999
Data columns (total 2 columns):
#   Column  Non-Null Count  Dtype
---  -
0   path    2000 non-null       object
1   label   2000 non-null       int64
dtypes: int64(1), object(1)
memory usage: 31.4+ KB
```

Grader function 2

In [9]:

```
1 def grader_df():
2     flag_shape = df_audio.shape==(2000,2)
3     flag_columns = all(df_audio.columns=='path', 'label'])
4     list_values = list(df_audio.label.value_counts())
5     flag_label = len(list_values)==10
6     flag_label2 = all([i==200 for i in list_values])
7     final_flag = flag_shape and flag_columns and flag_label and flag_la
8     return final_flag
9 grader_df()
```

Out[9]: True

In [10]:

```
1 from sklearn.utils import shuffle
2 df_audio = shuffle(df_audio, random_state=33)#don't change the random s
```

Train and Validation split

In [11]:

```
1 #split the data into train and validation and save in X_train, X_test,
2 #use stratify sampling
3 #use random state of 45
4 #use test size of 30%
5 X = df_audio['path']
6 y = df_audio['label']
7 X_train , X_test , y_train , y_test = train_test_split(X,y,test_size=0
8
```

Grader function 3

In [12]:

```
1 def grader_split():
2     flag_len = (len(X_train)==1400) and (len(X_test)==600) and (len(y_t
3     values_ytrain = list(y_train.value_counts())
4     flag_ytrain = (len(values_ytrain)==10) and (all([i==140 for i in va
5     values_ytest = list(y_test.value_counts())
6     flag_ytest = (len(values_ytest)==10) and (all([i==60 for i in value
7     final_flag = flag_len and flag_ytrain and flag_ytest
8     return final_flag
9 grader_split()
```

Out[12]: True

Preprocessing

All files are in the "WAV" format. We will read those raw data files using the librosa

```
In [13]: 1 sample_rate = 22050
2 def load_wav(x, get_duration=True):
3     '''This return the array values of audio with sampling rate of 22050
4     #loading the wav file with sampling rate of 22050
5     samples, sample_rate = librosa.load(x, sr=22050)
6     if get_duration:
7         duration = librosa.get_duration(samples, sample_rate)
8         return [samples, duration]
9     else:
10        return samples
```

```
In [14]: 1 #use load_wav function that was written above to get every wave.
2 #save it in X_train_processed and X_test_processed
3 # X_train_processed/X_test_processed should be dataframes with two columns
4 X_train_processed = []
5 X_test_processed = []
6
7 # X_train
8 for i in tqdm(range(len(X_train))):
9     processor = load_wav(X_train.iloc[i])
10    X_train_processed.append(processor)
11
12 #####
13 #X_test
14 for i in tqdm(range(len(X_test))):
15     processor = load_wav(X_test.iloc[i])
16    X_test_processed.append(processor)
17
18 100%|██████████| 1400/1400 [08:01<00:00, 2.91it/s]
19 100%|██████████| 600/600 [03:29<00:00, 2.86it/s]
```

```
In [15]: 1 train_1 = X_train_processed.copy()
2 test_1 = X_test_processed.copy()
```

```
In [16]: 1 X_train_processed[0]
```

```
Out[16]: [array([-0.00060508, -0.00061847, -0.00026167, ..., -0.00063194,
-0.00041468, 0.          ], dtype=float32), 0.3652607709750567]
```

```
In [18]: 1 new_X_train = pd.DataFrame(data={'raw_data':X_train_processed[i][0] for i in range(len(X_train_processed))})
2 new_X_test = pd.DataFrame(data={'raw_data':X_test_processed[i][0] for i in range(len(X_test_processed))})
```

In [19]: 1 new_X_test.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 600 entries, 0 to 599
Data columns (total 2 columns):
 #   Column      Non-Null Count  Dtype
---  -
 0   raw_data    600 non-null    object
 1   duration    600 non-null    float64
dtypes: float64(1), object(1)
memory usage: 9.5+ KB
```

In [20]: 1 new_X_train.head(2)

Out[20]:

	raw_data	duration
0	[-0.0006050776, -0.0006184709, -0.00026166602, ...]	0.365261
1	[0.00019366974, -0.0023843, -0.0058723832, -0....]	0.299773

In [21]: 1 new_X_test['raw_data'][0]

Out[21]: array([-4.9881153e-05, 6.9615439e-06, 5.9910067e-06, ...,
-3.3086265e-04, -2.4591145e-04, 0.0000000e+00], dtype=float32)

In [22]: 1 new_X_train.head(5)

Out[22]:

	raw_data	duration
0	[-0.0006050776, -0.0006184709, -0.00026166602, ...]	0.365261
1	[0.00019366974, -0.0023843, -0.0058723832, -0....]	0.299773
2	[-0.008321674, -0.013692323, -0.01608319, -0.0...]	0.246032
3	[0.0005516098, 0.00037881808, 0.00012904029, -...]	0.277415
4	[-0.00054838305, -0.000720711, -0.00075189554, ...]	0.472381

In [23]: 1 new_X_test.head(5)

Out[23]:


	raw_data	duration
0	[-4.9881153e-05, 6.961544e-06, 5.9910067e-06, ...]	0.400907
1	[0.0013453948, 0.0012860884, 0.00067011634, -0...]	0.477143
2	[-0.00095882645, -0.019986238, -0.02934129, -0...]	0.297778
3	[-9.786627e-06, -0.00016342092, -0.0004535091, ...]	0.447755
4	[0.011016414, 0.013351645, 0.013998778, 0.0138...]	0.537506

In [24]: 1 #new_X_train.to_csv('/gdrive/My Drive/Spoken_digit/X_train.csv', index=False)
2 #new_X_test.to_csv('/gdrive/My Drive/Spoken_digit/X_test.csv', index=False)

In [25]: 1 #new_X_train = pd.read_csv('/gdrive/My Drive/Spoken_digit/X_train.csv')
2 #new_X_test = pd.read_csv('/gdrive/My Drive/Spoken_digit/X_test.csv')
3


```
In [26]: 1 #plot the histogram of the duration for trian
          2 plt.hist(new_X_train['duration'])
```

```
Out[26]: (array([566., 713., 102., 15., 2., 0., 0., 0., 0., 2.]),
          array([0.14353741, 0.35746032, 0.57138322, 0.78530612, 0.99922902,
                  1.21315193, 1.42707483, 1.64099773, 1.85492063, 2.06884354,
                  2.28276644])),
          <a list of 10 Patch objects>)
```



```
In [27]: 1 #Plot for test data
          2 plt.hist(new_X_test['duration'])
```

```
Out[27]: (array([ 20.,  77., 137., 148., 110.,  54.,  37.,  11.,   3.,   3.]),
          array([0.14362812, 0.2158322 , 0.28803628, 0.36024036, 0.43244444,
                  0.50464853, 0.57685261, 0.64905669, 0.72126077, 0.79346485,
                  0.86566893])),
          <a list of 10 Patch objects>)
```



```
In [28]: 1 # refer - https://www.geeksforgeeks.org/numpy-percentile-in-python/
2 for i in range (0,101,10):
3     p = np.percentile(new_X_train['duration'], i)
4     print(str(i)+" Percentile: "+ str(p))
```

```
0 Percentile: 0.1435374149659864
10 Percentile: 0.2608934240362812
20 Percentile: 0.2977233560090703
30 Percentile: 0.3297777777777778
40 Percentile: 0.35663492063492064
50 Percentile: 0.389750566893424
60 Percentile: 0.41427664399092967
70 Percentile: 0.44360544217687076
80 Percentile: 0.4822312925170068
90 Percentile: 0.5535283446712018
100 Percentile: 2.282766439909297
```

```
In [29]: 1 # refer - https://www.geeksforgeeks.org/numpy-percentile-in-python/
2 for i in range (90,101,1):
3     p = np.percentile(new_X_train['duration'], i)
4     print(str(i)+" Percentile: "+ str(p))
```

```
90 Percentile: 0.5535283446712018
91 Percentile: 0.5659854875283448
92 Percentile: 0.5794503401360545
93 Percentile: 0.5938775510204082
94 Percentile: 0.6082149659863945
95 Percentile: 0.622421768707483
96 Percentile: 0.6424979591836734
97 Percentile: 0.6729219954648525
98 Percentile: 0.7120553287981859
99 Percentile: 0.8072766439909297
100 Percentile: 2.282766439909297
```

Grader function 4

```
In [30]: 1 X_train_processed = new_X_train
2 X_test_processed = new_X_test
```

```
In [31]: 1 def grader_processed():
2     flag_columns = (all(X_train_processed.columns==['raw_data', 'durati
3     flag_shape = (X_train_processed.shape ==(1400, 2)) and (X_test_proc
4     return flag_columns and flag_shape
5     grader_processed()
```

Out[31]: True

Based on our analysis 99 percentile values are less than 0.8sec so we will limit maximum length of X_train_processed and X_test_processed to 0.8 sec. It is similar to pad_sequence for a text dataset.

While loading the audio files, we are using sampling rate of 22050 so one sec will give array of length 22050. so, our maximum length is $0.8 \times 22050 = 17640$

In [32]: 1 max_length = 17640

In [33]:

```

1  ## as discussed above, Pad with Zero if length of sequence is less than
2  ## save in the X_train_pad_seq, X_test_pad_seq
3  ## also Create masking vector X_train_mask, X_test_mask
4
5  ## all the X_train_pad_seq, X_test_pad_seq, X_train_mask, X_test_mask w
6  #X_train padding Sequences
7  X_train_pad_seq = []
8  for i in tqdm(range(len(new_X_train))):
9      sequences = [new_X_train['raw_data'][i]]
10     aaa = tf.keras.preprocessing.sequence.pad_sequences(
11         sequences, maxlen=max_length, dtype='float', padding='post', trunc
12         value=100)
13     X_train_pad_seq.extend(aaa)
14
15     #####
16
17     #X_test padding Sequences
18     X_test_pad_seq = []
19     for i in tqdm(range(len(new_X_test))):
20         sequences = [new_X_test['raw_data'][i]]
21         aaa = tf.keras.preprocessing.sequence.pad_sequences(
22             sequences, maxlen=max_length, dtype='float', padding='post', trunc
23             value=100)
24         X_test_pad_seq.extend(aaa)
25
100%|██████████| 1400/1400 [00:00<00:00, 8257.82it/s]
100%|██████████| 600/600 [00:00<00:00, 8749.22it/s]

```

In [34]:

```

1  X_train_pad_seq = np.array(X_train_pad_seq)
2  X_test_pad_seq = np.array(X_test_pad_seq)
3  print('shape of X_train padding', X_train_pad_seq.shape)
4  print('shape of X_test padding', X_test_pad_seq.shape)

```

shape of X_train padding (1400, 17640)

shape of X_test padding (600, 17640)


```
In [35]: 1 # X_train_mask
2 X_train_mask = []
3 for i in tqdm(range(len(new_X_train))):
4     X_train_mask_replace = np.where(X_train_pad_seq[i]!=100.0, 1, X_train
5     X_train_mask_replace = np.where(X_train_mask_replace==100.0, 0, X_tra
6     X_train_mask.append(X_train_mask_replace)
7
8 #####
9
10 X_test_mask = []
11 for i in tqdm(range(len(new_X_test))):
12     X_test_mask_replace = np.where(X_test_pad_seq[i]!=100.0, 1, X_train_p
13     X_test_mask_replace = np.where(X_test_mask_replace==100.0, 0, X_test_
14     X_test_mask.append(X_test_mask_replace)

100%|██████████| 1400/1400 [00:00<00:00, 11055.14it/s]
100%|██████████| 600/600 [00:00<00:00, 11506.64it/s]
```

```
In [36]: 1 X_train_mask[:2]
```

```
Out[36]: [array([1., 1., 1., ..., 0., 0., 0.]), array([1., 1., 1., ..., 0., 0.,
0.])]
```

```
In [37]: 1 X_train_mask = np.array(X_train_mask).astype('bool')
2 X_test_mask = np.array(X_test_mask).astype('bool')
```

```
In [38]: 1 X_train_mask[:2]
```

```
Out[38]: array([[ True,  True,  True, ..., False, False, False],
[ True,  True,  True, ..., False, False, False]])
```

Grader function 5

```
In [39]: 1 def grader_padoutput():
2     flag_padshape = (X_train_pad_seq.shape==(1400, 17640)) and (X_test_
3     #print(flag_padshape)
4     flag_maskshape = (X_train_mask.shape==(1400, 17640)) and (X_test_ma
5     #print(flag_maskshape)
6     flag_dtype = (X_train_mask.dtype==bool) and (X_test_mask.dtype==boo
7     #print(flag_dtype)
8     return flag_padshape and flag_maskshape and flag_dtype
9     grader_padoutput()
```

```
Out[39]: True
```

1. Giving Raw data directly.

Now we have

Train data: X_train_pad_seq, X_train_mask and y_train

Test data: X_test_pad_seq, X_test_mask and y_test

We will create a LSTM model which takes this input.

Task:

1. Create an LSTM network which takes "X_train_pad_seq" as input, "X_train_mask" as mask input. You can use any number of LSTM cells. Please read LSTM documentation(https://www.tensorflow.org/api_docs/python/tf/keras/layers/LSTM) in tensorflow to know more about mask and also https://www.tensorflow.org/guide/keras/masking_and_padding
2. Get the final output of the LSTM and give it to Dense layer of

```
In [40]: 1 from tensorflow.keras.layers import Input, LSTM, Dense,Masking,Dropout
2 from tensorflow.keras.models import Model
3 import tensorflow as tf
4 from sklearn.metrics import confusion_matrix, f1_score, precision_score
5 from keras.regularizers import l2,l1_l2,l1
6 from tensorflow.keras.callbacks import ModelCheckpoint, EarlyStopping,
7 import datetime
8 from tensorflow.keras.layers import AveragePooling1D,GlobalAveragePool
9 import random as rn
10
```

In []:

```

1  ## as discussed above, please write the LSTM
2  time_steps = 17640
3  n_features = 1
4  #this is input words. Sequence of words represented as integers
5  input_padding = tf.keras.layers.Input(shape=(time_steps,n_features), na
6
7  #mask vector if you are padding anything
8  input_mask = tf.keras.layers.Input(shape=(time_steps), name="input_Mask
9
10
11 lstm = LSTM(25)(input_padding)
12 x = Dense(50, activation='relu',kernel_initializer=tf.keras.initializer
13 x = Dropout(0.2)(x)
14 output = Dense(10, activation = 'softmax')(x)
15
16 model = Model(inputs=[input_mask,input_padding],outputs=output)
17 model.summary()

```

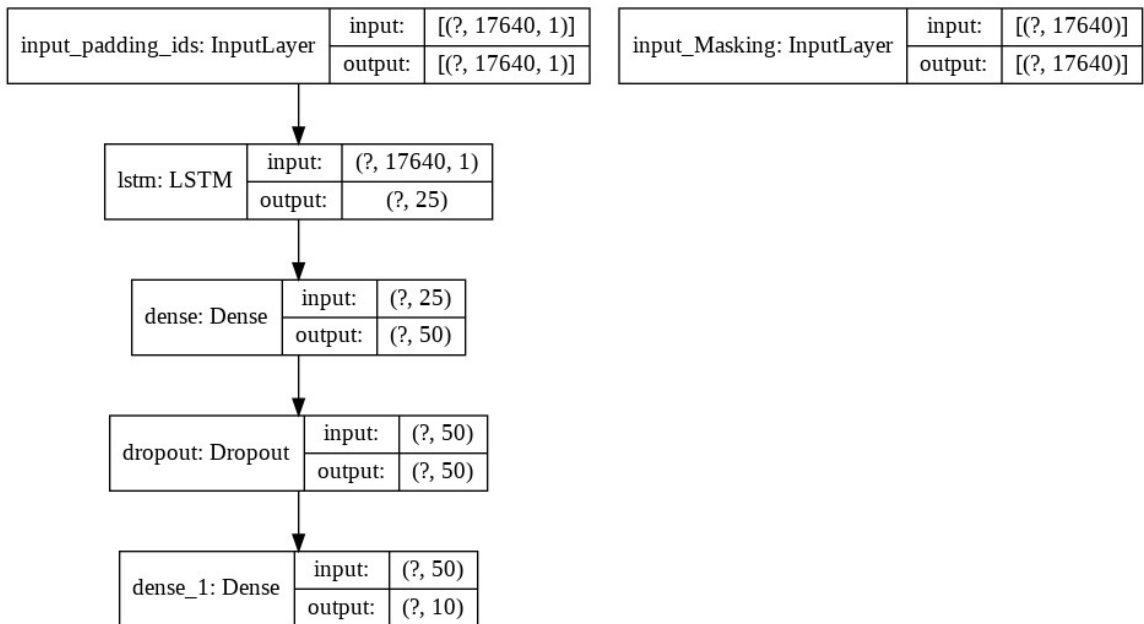
Model: "functional_13"

Layer (type)	Output Shape	Param #	Connected to
input_padding_ids (InputLayer)	[(None, 17640, 1)]	0	
lstm_11 (LSTM)	(None, 25)	2700	input_padding_ids[0][0]
dense_14 (Dense)	(None, 50)	1300	lstm_11
dropout_7 (Dropout)	(None, 50)	0	dense_14
input_Masking (InputLayer)	[(None, 17640)]	0	
dense_15 (Dense)	(None, 10)	510	dropout_7

Total params: 4,510
 Trainable params: 4,510
 Non-trainable params: 0

```
In [ ]: 1 dot_img_file = '/tmp/model_1.png'
2 tf.keras.utils.plot_model(model, to_file=dot_img_file, show_shapes=True)
```

Out[28]:



```
In [ ]: 1 ACCURACY_THRESHOLD = 0.1
2 class Metrics(tf.keras.callbacks.Callback):
3     def __init__(self, validation):
4         super(Metrics, self).__init__()
5         self.validation_data = validation
6
7     def on_train_begin(self, logs={}):
8         self.val_f1s = []
9
10
11     def on_epoch_end(self, epoch, logs={}):
12
13         val_predict = self.model.predict(self.validation_data[0])
14         val_predict = np.argmax(val_predict, axis=1)
15
16         val_targ = self.validation_data[1]
17
18         _val_f1 = f1_score(val_targ, val_predict, average='micro')
19         self.val_f1s.append(_val_f1)
20
21         print(' - val_f1: %f ' %(_val_f1))
22
23         if (_val_f1 > ACCURACY_THRESHOLD):
24             print("\nReached %2.2f%% accuracy, so stopping training!!" %(ACCU
25                 self.model.stop_training = True
26
27
28 metrics = Metrics((X_train_pad_seq, y_train))
```

```
In [ ]: 1 # Call back
2 earlystop = EarlyStopping(monitor='val_accuracy', min_delta=0, patience
3 filepath="/gdrive/My Drive/model_save/best_model-{epoch:02d}.h5"
4 checkpoint = ModelCheckpoint(filepath=filepath, monitor='val_accuracy',
5 # TensorBoard Creation
6 %load_ext tensorboard
7 folder_name = datetime.datetime.now().strftime("%Y%m%d-%H%M%S")
8
9 # Create Log folder - TensorBoard
10 log_dir="/gdrive/My Drive/logs/fit/" + folder_name
11 tensorboard_callback = TensorBoard(log_dir=log_dir, histogram_freq=1, wri
```

The tensorboard extension is already loaded. To reload it, use:
%reload_ext tensorboard

```
In [ ]: 1 folder_name
```

```
Out[103]: '20201112-144741'
```

Train data: X_train_pad_seq, X_train_mask and y_train

Test data: X_test_pad_seq, X_test_mask and y_test

```
In [ ]: 1 model.compile(loss='sparse_categorical_crossentropy', optimizer= 'adam'
2 model.fit(x = X_train_pad_seq, y=y_train, epochs=40, verbose=1, batch_si
3         callbacks = [checkpoint, tensorboard_callback, earlystop, metrics
```

Epoch 1/40

2/22 [=>.....] - ETA: 41s - loss: 2.2067 - accurac
y: 0.1797WARNING:tensorflow:Callbacks method `on_train_batch_end` is slow
compared to the batch time (batch time: 0.9592s vs `on_train_batch_end` ti
me: 3.1906s). Check your callbacks.

22/22 [=====] - ETA: 0s - loss: 2.1795 - accurac
y: 0.1607

Epoch 00001: val_accuracy improved from 0.16000 to 0.17286, saving model t
o /gdrive/My Drive/model_save/best_model-01.h5
- val_f1: 0.172857


Reached 10.00% accuracy, so stopping training!!


22/22 [=====] - 43s 2s/step - loss: 2.1795 - accu
racy: 0.1607 - val_loss: 2.1571 - val_accuracy: 0.1729

```
Out[99]: <tensorflow.python.keras.callbacks.History at 0x7f6e58218fd0>
```

```
In [ ]: 1 #Model 1 - results
2 os.chdir('/gdrive/My Drive/')
3 %tensorboard --logdir logs/fit/
4
```

Output hidden; open in <https://colab.research.google.com> (<https://colab.research.google.com>) to view.





2. Converting into spectrogram and giving spectrogram data as input

We can use librosa to convert raw data into spectrogram. A spectrogram shows the features in a two-dimensional representation with the

intensity of a frequency at a point in time i.e we are converting Time domain to frequency domain. you can read more about this in <https://pnsn.org/spectrograms/what-is-a-spectrogram>

```
In [41]: 1 def convert_to_spectrogram(raw_data):
2         '''converting to spectrogram'''
3
4         spectrum = librosa.feature.melspectrogram(y=raw_data, sr=sample_rate)
5         logmel_spectrum = librosa.power_to_db(S=spectrum, ref=np.max)
6         #print(logmel_spectrum.shape)
7         return logmel_spectrum
```

```
In [42]: 1 ##use convert_to_spectrogram and convert every raw sequence in X_train_
2         ## save those all in the X_train_spectrogram and X_test_spectrogram ( T
3         ##Train data: X_train_pad_seq, X_train_mask and y_train
4         ##Test data: X_test_pad_seq, X_test_mask and y_test
5
6
7         X_train_spectrogram = []
8         for i in tqdm(range(len(X_train_pad_seq))):
9             aaa = convert_to_spectrogram(X_train_pad_seq[i])
10            X_train_spectrogram.append(aaa)
11
12
13            #####
14
15            #X_test padding Sequences
16            X_test_spectrogram = []
17            for i in tqdm(range(len(X_test_pad_seq))):
18                bbb = convert_to_spectrogram(X_test_pad_seq[i])
19                X_test_spectrogram.append(bbb)
20
21
```

```
100%|██████████| 1400/1400 [00:08<00:00, 168.63it/s]
100%|██████████| 600/600 [00:03<00:00, 170.37it/s]
```

```
In [43]: 1 X_train_spectrogram = np.array(X_train_spectrogram)
          2 X_test_spectrogram = np.array(X_test_spectrogram)
          3 print('shape of X_train spectrogram', X_train_spectrogram.shape)
          4 print('shape of X_test spectrogram', X_test_spectrogram.shape)
```

shape of X_train spectrogram (1400, 64, 35)

shape of X_test spectrogram (600, 64, 35)

Grader function 6

```
In [44]: 1 def grader_spectrogram():
          2     flag_shape = (X_train_spectrogram.shape==(1400,64, 35)) and (X_test
          3     return flag_shape
          4     grader_spectrogram()
```

Out[44]: True

Now we have

Train data: X_train_spectrogram and y_train

Test data: X_test_spectrogram and y_test

We will create a LSTM model which takes this input.

Task:

1. Create an LSTM network which takes "X_train_spectrogram" as input and has to return output at every time step.
2. Average the output of every time step and give this to the Dense layer of any size.
(ex: Output from LSTM will be (#., time_steps, features) average the output of every time step i.e, you should get (#.,time_steps) and then pass to dense layer)
3. give the above output to Dense layer of size 10(output layer) and train the network with sparse categorical cross entropy.
4. Use tensorboard to plot the graphs of loss and metric(use micro F1 score as metric) and histograms of gradients.
5. make sure that it won't overfit.
6. You are free to include any regularization

In [51]:

```

1 tf.keras.backend.clear_session()
2
3 ## Set the random seed values to regenerate the model.
4 np.random.seed(0)
5 rn.seed(0)
6
7
8 ## as discussed above, please write the LSTM
9 time_steps = 64
10 n_features = 35
11 #this is input words. Sequence of words represented as integers
12 input_padding = tf.keras.layers.Input(shape=(time_steps,n_features), na
13 lstm = LSTM(2000,return_sequences=True)(input_padding)
14 global_average = GlobalAveragePooling1D(data_format='channels_first')(
15 #res = tf.reduce_mean(global_average , axis = 1, keepdims = True)
16 x = Dense(1000, activation = 'relu',kernel_initializer=tf.keras.initial
17 x = Dropout(rate=0.5)(x)
18 output = Dense(10, activation = 'softmax')(x)
19
20 model = Model(inputs=input_padding,outputs=output)
21 model.summary()

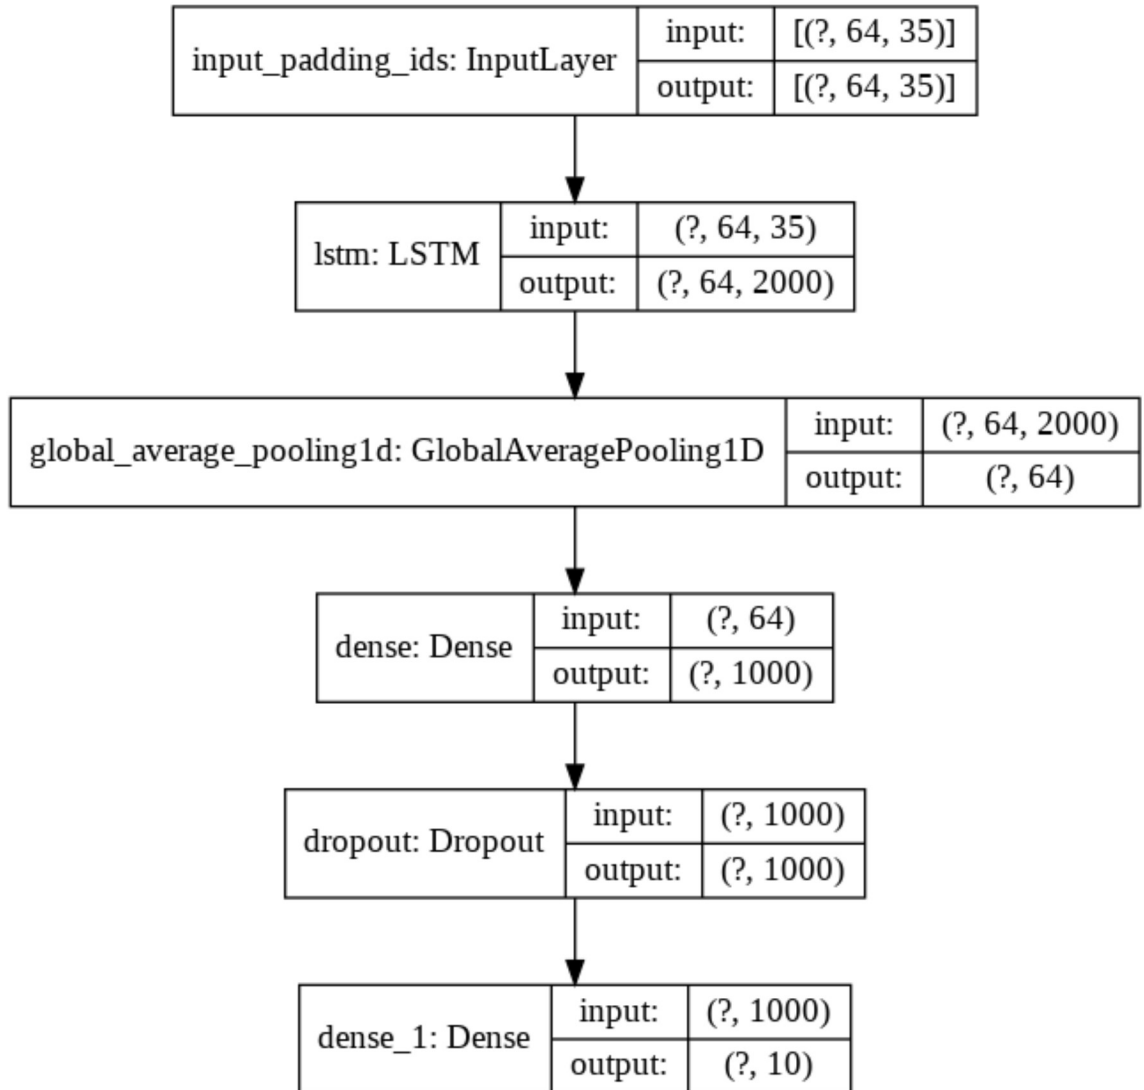
```

Model: "functional_1"

Layer (type)	Output Shape	Param #
=====		
input_padding_ids (InputLaye	[(None, 64, 35)]	0
=====		
lstm (LSTM)	(None, 64, 2000)	16288000
=====		
global_average_pooling1d (Gl	(None, 64)	0
=====		
dense (Dense)	(None, 1000)	65000
=====		
dropout (Dropout)	(None, 1000)	0
=====		
dense_1 (Dense)	(None, 10)	10010
=====		
Total params: 16,363,010		
Trainable params: 16,363,010		
Non-trainable params: 0		
=====		


```
In [52]: 1 dot_img_file = '/tmp/model_1.png'
2         tf.keras.utils.plot_model(model, to_file=dot_img_file, show_shapes=True)
```

Out[52]:



In [53]:

```
1 ACCURACY_THRESHOLD = 0.8
2 class Metrics(tf.keras.callbacks.Callback):
3     def __init__(self, validation):
4         super(Metrics, self).__init__()
5         self.validation_data = validation
6
7     def on_train_begin(self, logs={}):
8         self.val_f1s_score = []
9         self.f1_score_best = 0
10        self.epoch_value = 1
11
12
13    def on_epoch_end(self, epoch, logs={}):
14
15
16
17        val_predict = self.model.predict(self.validation_data[0])
18        val_predict = np.argmax(val_predict, axis=1)
19
20        val_targ = self.validation_data[1]
21
22        _val_f1 = f1_score(val_targ, val_predict, average='micro')
23        self.val_f1s_score.append(_val_f1)
24
25        print(' - val_f1: %f ' %(_val_f1))
26
27
28        if _val_f1 > self.f1_score_best:
29            print('F1_score improved from '+str(self.f1_score_best ) + ' to '
30                  self.f1_score_best = _val_f1
31        else:
32            print('Model not improved, still best f1_score reamins '+str(sel
33
34        self.epoch_value = self.epoch_value +1
35        if (_val_f1 >= ACCURACY_THRESHOLD):
36            print("\nReached %2.2f%% accuracy, so stopping training!!" %(ACCU
37                  self.model.stop_training = True
38            return
39
40 metrics = Metrics((X_test_spectrogram, y_test))
```

In [54]:

```

1  # Call back
2
3  # Learning rate scheduler
4  def my_learning_rate(epoch, lrate):
5      if epoch <=500:
6          print('Learning rate changed to 0.0001')
7          lrate = 0.0001
8
9      else:
10         print('Learning rate changed to 0.00001')
11         lrate = 0.00001
12
13     return lrate
14
15 lrs = LearningRateScheduler(my_learning_rate)
16 red_learn = tf.keras.callbacks.ReduceLROnPlateau(
17     monitor="val_accuracy",
18     factor=0.1,
19     patience=15,
20     verbose=0,
21     mode="auto",
22     min_delta=0.0001,
23     cooldown=0,
24     min_lr=0
25 )
26 earlystop = EarlyStopping(monitor='val_accuracy', min_delta=0, patience
27 filepath="/gdrive/My Drive/model_save/best_model2-{epoch:02d}.h5"
28 checkpoint = ModelCheckpoint(filepath=filepath, monitor='val_accuracy',
29 # TensorBoard Creation
30 %load_ext tensorboard
31 folder_name = datetime.datetime.now().strftime("%Y%m%d-%H%M%S")
32
33 # Create Log folder - TensorBoard
34 log_dir="/gdrive/My Drive/logs/fit/" + folder_name
35 tensorboard_callback =TensorBoard(log_dir=log_dir,histogram_freq=1, wri

```

The tensorboard extension is already loaded. To reload it, use:
 %reload_ext tensorboard

In [55]:

```
1 folder_name
```

Out[55]: '20201122-030151'

```
In [56]: 1 optim = tf.keras.optimizers.Adam(learning_rate=0.0001)
2 model.compile(loss='sparse_categorical_crossentropy', optimizer= optim,
3 model.fit(x = X_train_spectrogram, y=y_train, epochs=1000,verbose=1, v
4         callbacks =[earlystop,tensorboard_callback,metrics])
```

Epoch 1/1000

2/44 [>.....] - ETA: 4s - loss: 2.3051 - accuracy: 0.1250WARNING:tensorflow:Callbacks method `on_train_batch_end` is slow compared to the batch time (batch time: 0.0419s vs `on_train_batch_end` time: 0.1813s). Check your callbacks.

44/44 [=====] - ETA: 0s - loss: 2.2915 - accuracy: 0.1121 - val_f1: 0.143333

F1_score improved from 0 to 0.14333333333333334 Epoch value 1

44/44 [=====] - 7s 160ms/step - loss: 2.2915 - accuracy: 0.1121 - val_loss: 2.2769 - val_accuracy: 0.1433

Epoch 2/1000

44/44 [=====] - ETA: 0s - loss: 2.2305 - accuracy: 0.1464 - val_f1: 0.145000

F1_score improved from 0.14333333333333334 to 0.145 Epoch value 2

44/44 [=====] - 6s 145ms/step - loss: 2.2305 - accuracy: 0.1464 - val_loss: 2.2373 - val_accuracy: 0.1450


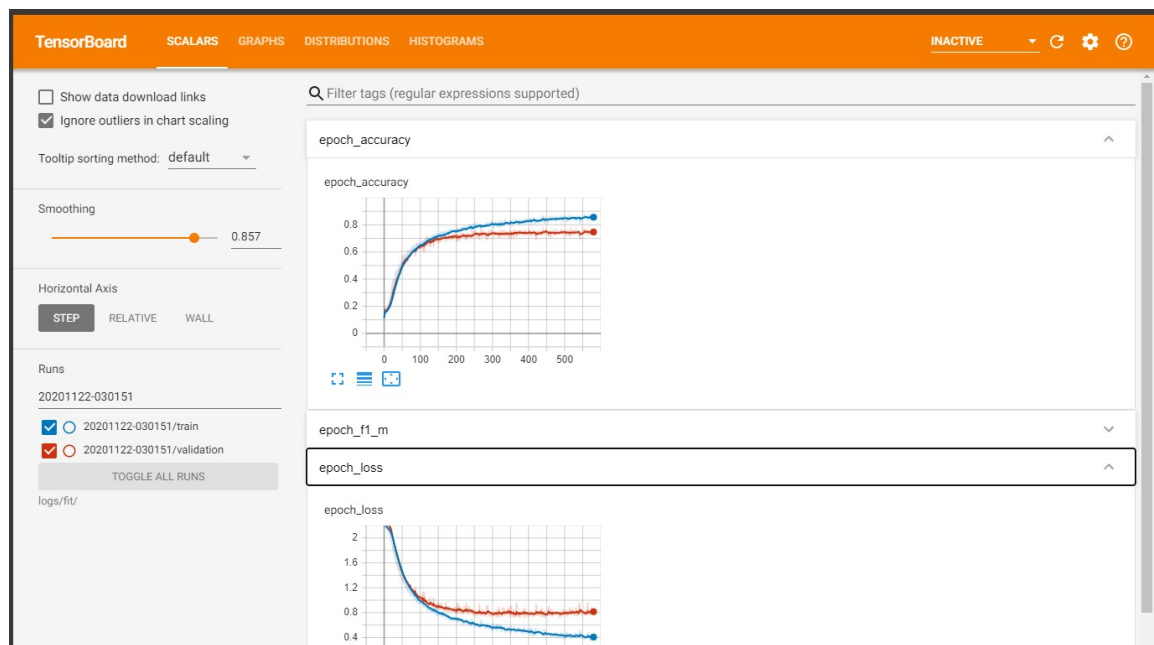
Epoch 3/1000


44/44 [=====] - ETA: 0s - loss: 2.1848 - accuracy: 0.1793 - val_f1: 0.171667

F1_score improved from 0.145 to 0.17166666666666666 Epoch value 3

```
In [58]: 1 #Model 1 - results
2 os.chdir('/gdrive/My Drive/')
3 %tensorboard --logdir logs/fit/
```

Output hidden; open in <https://colab.research.google.com> (<https://colab.research.google.com>) to view.





3. data augmentation

Till now we have done with 2000 samples only. It is very less data. We are giving the process of generating augmented data below.

There are two types of augmentation:

1. time stretching - Time stretching either increases or decreases the length of the file. For time stretching we move the file 30% faster or slower
2. pitch shifting - pitch shifting moves the frequencies higher or lower. For pitch shifting we shift up or down one half-step.

```
In [59]: 1 ## generating augmented data.
          2 def generate_augmented_data(file_path):
          3     augmented_data = []
          4     samples = load_wav(file_path, get_duration=False)
          5     for time_value in [0.7, 1, 1.3]:
          6         for pitch_value in [-1, 0, 1]:
          7             time_stretch_data = librosa.effects.time_stretch(samples, r
          8             final_data = librosa.effects.pitch_shift(time_stretch_data,
          9             augmented_data.append(final_data)
          10     return augmented_data
```

```
In [60]: 1 temp_path = df_audio.iloc[0].path
          2 aug_temp = generate_augmented_data(temp_path)
```

```
In [61]: 1 aug_temp[8]
```

```
Out[61]: array([-0.00034318, -0.00020249, -0.00010748, ...,  0.0001497 ,
                0.0002406 ,  0.00029225], dtype=float32)
```

```
In [62]: 1 len(aug_temp)
```

```
Out[62]: 9
```

```
In [63]: 1 df_audio.iloc[0].label
```

```
Out[63]: 7
```

As discussed above, for one data point, we will get 9 augmented data points.

Split data into train and test (80-20 split)

We have 2000 data points(1600 train points, 400 test points)

Do augmentation only on train data, after augmentation we will get 14400 train points.

```
In [64]: 1 X = df_audio['path']
          2 y = df_audio['label']
          3 X_train , X_test , y_train , y_test = train_test_split(X,y,test_size=0
          4
```

```
In [65]: 1 print(X_train.iloc[0])
          2 print(y_train.iloc[0])

/gdrive/My Drive/recordings/8_jackson_35.wav
8
```

```
In [66]: 1 X_train.head()
```

```
Out[66]: 597      /gdrive/My Drive/recordings/8_jackson_35.wav
          126      /gdrive/My Drive/recordings/2_nicolas_13.wav
          1763     /gdrive/My Drive/recordings/7_yeweler_40.wav
          302      /gdrive/My Drive/recordings/3_theo_2.wav
          1493     /gdrive/My Drive/recordings/9_yeweler_27.wav
          Name: path, dtype: object
```

```
In [67]: 1 X_train_spectrogram_augument = []
          2 y_train_augument = []
          3 y1 = 0
          4 for i in tqdm(range(len(X_train))):
          5     aaa = generate_augmented_data(X_train.iloc[i])
          6     for k in range(len(aaa)):
          7         value = aaa[k]
          8         label = y_train.iloc[y1]
          9         X_train_spectrogram_augument.append(value)
          10        y_train_augument.append(label)
          11        y1+=1
          12
          13 # Augumentation only to train data
          14 X_test_spectrogram_augument = []
          15 y_test_augument = []
          16 for i in tqdm(range(len(X_test))):
          17     processor = load_wav(X_test.iloc[i], get_duration=False)
          18     X_test_spectrogram_augument.append(processor)
          19     y_test_augument.append(y_test.iloc[i])
          20
          21

100%|██████████| 1600/1600 [13:46<00:00, 1.94it/s]
100%|██████████| 400/400 [02:15<00:00, 2.95it/s]
```

```
In [68]: 1 new_X_train_augument = pd.DataFrame({'raw_input':X_train_spectrogram_augument,
          2 new_X_test_augument = pd.DataFrame({'raw_input':X_test_spectrogram_augument,
```

```
In [69]: 1 new_X_train_augument.head(2)
```

```
Out[69]:
```

	raw_input	label
0	[0.005341887, 0.0069970735, 0.007084101, 0.007...	8
1	[0.006056736, 0.0074520707, 0.007774388, 0.007...	8

In [70]: 1 new_X_train_augument.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 14400 entries, 0 to 14399
Data columns (total 2 columns):
#   Column      Non-Null Count  Dtype
---  -
0   raw_input    14400 non-null  object
1   label        14400 non-null  int64
dtypes: int64(1), object(1)
memory usage: 225.1+ KB
```

In [71]: 1 new_X_test_augument.info()

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 400 entries, 0 to 399
Data columns (total 2 columns):
#   Column      Non-Null Count  Dtype
---  -
0   raw_input    400 non-null   object
1   label        400 non-null   int64
dtypes: int64(1), object(1)
memory usage: 6.4+ KB
```

In [72]: 1 X_train_processed = new_X_train_augument
2 X_test_processed = new_X_test_augument

In [73]:

```
1 max_length = 17640
2
3 ## as discussed above, Pad with Zero if length of sequence is less than
4 ## save in the X_train_pad_seq, X_test_pad_seq
5 ## also Create masking vector X_train_mask, X_test_mask
6
7 ## all the X_train_pad_seq, X_test_pad_seq, X_train_mask, X_test_mask w
8 #X_train padding Sequences
9 X_train_pad_seq = []
10 for i in tqdm(range(len(new_X_train_augument))):
11     sequences = [new_X_train_augument['raw_input'][i]]
12     aaa = tf.keras.preprocessing.sequence.pad_sequences(
13         sequences, maxlen=max_length, dtype='float', padding='post', trun
14         value=100)
15     X_train_pad_seq.extend(aaa)
16
17 #####
18
19 #X_test padding Sequences
20 X_test_pad_seq = []
21 for i in tqdm(range(len(new_X_test_augument))):
22     sequences = [new_X_test_augument['raw_input'][i]]
23     aaa = tf.keras.preprocessing.sequence.pad_sequences(
24         sequences, maxlen=max_length, dtype='float', padding='post', trun
25         value=100)
26     X_test_pad_seq.extend(aaa)
27
```

```
100%|██████████| 14400/14400 [00:01<00:00, 8781.28it/s]
100%|██████████| 400/400 [00:00<00:00, 15583.23it/s]
```

```
In [74]: 1 X_train_pad_seq = np.array(X_train_pad_seq)
2 X_test_pad_seq = np.array(X_test_pad_seq)
3 print('shape of X_train padding', X_train_pad_seq.shape)
4 print('shape of X_test padding', X_test_pad_seq.shape)
```

shape of X_train padding (14400, 17640)
 shape of X_test padding (400, 17640)

```
In [75]: 1 # X_train_mask
2 X_train_mask = []
3 for i in tqdm(range(len(new_X_train_augument))):
4     X_train_mask_replace = np.where(X_train_pad_seq[i]!=100.0, 1, X_train
5     X_train_mask_replace = np.where(X_train_mask_replace==100.0, 0, X_tra
6     X_train_mask.append(X_train_mask_replace)
7
8 #####
9
10 X_test_mask = []
11 for i in tqdm(range(len(new_X_test_augument))):
12     X_test_mask_replace = np.where(X_test_pad_seq[i]!=100.0, 1, X_train_p
13     X_test_mask_replace = np.where(X_test_mask_replace==100.0, 0, X_test_
14     X_test_mask.append(X_test_mask_replace)
```

100%|██████████| 14400/14400 [00:01<00:00, 9347.69it/s]

100%|██████████| 400/400 [00:00<00:00, 10266.38it/s]

```
In [76]: 1 X_train_mask[:2]
```

```
Out[76]: [array([1., 1., 1., ..., 0., 0., 0.]), array([1., 1., 1., ..., 0., 0.,
0.])]
```

```
In [77]: 1 X_train_mask = np.array(X_train_mask).astype('bool')
2 X_test_mask = np.array(X_test_mask).astype('bool')
```

```
In [78]: 1 X_train_mask[:2]
```

```
Out[78]: array([[ True,  True,  True, ..., False, False, False],
[ True,  True,  True, ..., False, False, False]])
```

```
In [79]: 1 from tensorflow.keras.layers import Input, LSTM, Dense,Masking,Dropout
2 from tensorflow.keras.models import Model
3 import tensorflow as tf
4 from sklearn.metrics import confusion_matrix, f1_score, precision_score
5 from keras.regularizers import l2,l1_l2,l1
6 from tensorflow.keras.callbacks import ModelCheckpoint, EarlyStopping,
7 import datetime
8 from tensorflow.keras.layers import AveragePooling1D,GlobalAveragePool
9 import random as rn
10
```


In []:

```

1  ## as discussed above, please write the LSTM
2  time_steps = 17640
3  n_features = 1
4  #this is input words. Sequence of words represented as integers
5  input_padding = tf.keras.layers.Input(shape=(time_steps,n_features), na
6
7  #mask vector if you are padding anything
8  input_mask = tf.keras.layers.Input(shape=(time_steps), name="input_Mask
9
10
11 lstm = LSTM(128)(input_padding)
12 x = Dense(64, activation='relu',kernel_initializer=tf.keras.initializer
13 x = Dropout(0.2)(x)
14 output = Dense(10, activation = 'softmax')(x)
15
16 model = Model(inputs=[input_padding],outputs=output)
17 model.summary()

```

Model: "functional_9"

Layer (type)	Output Shape	Param #
=====		
input_padding_ids (InputLaye	[(None, 17640, 1)]	0
=====		
lstm_5 (LSTM)	(None, 128)	66560
=====		
dense_8 (Dense)	(None, 64)	8256
=====		
dropout_3 (Dropout)	(None, 64)	0
=====		
dense_9 (Dense)	(None, 10)	650
=====		
Total params: 75,466		
Trainable params: 75,466		
Non-trainable params: 0		

In [85]:

```

1  y_train_fit = new_X_train_augument['label']
2  y_test_fit = new_X_test_augument['label']

```

```

In [ ]: 1 ACCURACY_THRESHOLD = 0.1
        2 class Metrics(tf.keras.callbacks.Callback):
        3     def __init__(self, validation):
        4         super(Metrics, self).__init__()
        5         self.validation_data = validation
        6
        7     def on_train_begin(self, logs={}):
        8         self.val_f1s = []
        9
        10
        11     def on_epoch_end(self, epoch, logs={}):
        12
        13         val_predict = self.model.predict(self.validation_data[0])
        14         val_predict = np.argmax(val_predict, axis=1)
        15
        16         val_targ = self.validation_data[1]
        17
        18         _val_f1 = f1_score(val_targ, val_predict, average='micro')
        19         self.val_f1s.append(_val_f1)
        20
        21         print(' - val_f1: %f ' %(_val_f1))
        22
        23         if (_val_f1 > ACCURACY_THRESHOLD):
        24             print("\nReached %2.2f%% accuracy, so stopping training!!" %(ACCU
        25                 self.model.stop_training = True
        26
        27
        28 metrics = Metrics((X_test_pad_seq, y_test_fit))

```

```

In [ ]: 1 # Call back
        2 earlystop = EarlyStopping(monitor='val_accuracy', min_delta=0, patience
        3 filepath="/gdrive/My Drive/model_save/best_model-{epoch:02d}.h5"
        4 checkpoint = ModelCheckpoint(filepath=filepath, monitor='val_accuracy',
        5 # TensorBoard Creation
        6 %load_ext tensorboard
        7 folder_name = datetime.datetime.now().strftime("%Y%m%d-%H%M%S")
        8
        9 # Create Log folder - TensorBoard
        10 log_dir="/gdrive/My Drive/logs/fit/" + folder_name
        11 tensorboard_callback =TensorBoard(log_dir=log_dir,histogram_freq=1, wri

```

The tensorboard extension is already loaded. To reload it, use:
%reload_ext tensorboard

```

In [ ]: 1 folder_name

```

Out[63]: '20201120-160801'

```

In [ ]: 1

```

```
In [ ]: 1 model.compile(loss='sparse_categorical_crossentropy', optimizer= 'adam'
2 model.fit(x = X_train_pad_seq, y=y_train_fit, epochs=40,verbose=1, val
3         callbacks =[checkpoint,tensorboard_callback,earlystop,metrics
```

Epoch 1/40

1/450 [.....] - ETA: 0s - loss: 2.4686 - accuracy: 0.0938
 WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/tensorflow/python/ops/summary_ops_v2.py:1277: stop (from tensorflow.python.eager.profiler) is deprecated and will be removed after 2020-07-01.

Instructions for updating:

use `tf.profiler.experimental.stop` instead.

2/450 [.....] - ETA: 14:12 - loss: 2.4677 - accuracy: 0.0625
 WARNING:tensorflow:Callbacks method `on_train_batch_end` is slow compared to the batch time (batch time: 1.1473s vs `on_train_batch_end` time: 2.6599s). Check your callbacks.

450/450 [=====] - ETA: 0s - loss: 2.3075 - accuracy: 0.1081

Epoch 00001: val_accuracy improved from -inf to 0.10000, saving model to /gdrive/My Drive/model_save/best_model-01.h5

- val_f1: 0.100000

Reached 10.00% accuracy, so stopping training!!

450/450 [=====] - 337s 750ms/step - loss: 2.3075
 - accuracy: 0.1081 - val_loss: 2.3090 - val_accuracy: 0.1000

Out[64]: <tensorflow.python.keras.callbacks.History at 0x7f788a72da90>

MODEL - 4

```
In [80]: 1 ##use convert_to_spectrogram and convert every raw sequence in X_train_
2 ## save those all in the X_train_spectrogram and X_test_spectrogram ( T
3 ##Train data: X_train_pad_seq, X_train_mask and y_train
4 #Test data: X_test_pad_seq, X_test_mask and y_test
5
6
7 X_train_spectrogram = []
8 for i in tqdm(range(len(X_train_pad_seq))):
9     aaa = convert_to_spectrogram(X_train_pad_seq[i])
10    X_train_spectrogram.append(aaa)
11
12
13 #####
14
15 #X_test padding Sequences
16 X_test_spectrogram = []
17 for i in tqdm(range(len(X_test_pad_seq))):
18     bbb = convert_to_spectrogram(X_test_pad_seq[i])
19     X_test_spectrogram.append(bbb)
20
21
```

100%|██████████| 14400/14400 [01:29<00:00, 160.27it/s]

100%|██████████| 400/400 [00:03<00:00, 127.82it/s]

```
In [81]: 1 X_train_spectrogram = np.array(X_train_spectrogram)
2 X_test_spectrogram = np.array(X_test_spectrogram)
3 print('shape of X_train spectrogram', X_train_spectrogram.shape)
4 print('shape of X_test spectrogram', X_test_spectrogram.shape)
```

shape of X_train spectrogram (14400, 64, 35)
 shape of X_test spectrogram (400, 64, 35)

```
In [82]: 1 tf.keras.backend.clear_session()
2
3 ## Set the random seed values to regenerate the model.
4 np.random.seed(0)
5 rn.seed(0)
6
7
8 ## as discussed above, please write the LSTM
9 time_steps = 64
10 n_features = 35
11 #this is input words. Sequence of words represented as integers
12 input_padding = tf.keras.layers.Input(shape=(time_steps,n_features), name='input')
13 lstm = LSTM(1500,return_sequences=True)(input_padding)
14 global_average = GlobalAveragePooling1D(data_format='channels_first')(lstm)
15 #res = tf.reduce_mean(global_average , axis = 1, keepdims = True)
16 x = Dense(1000, activation = 'relu',kernel_initializer=tf.keras.initializers.glorot_uniform)(global_average)
17 x = Dropout(rate=0.4)(x)
18 output = Dense(10, activation = 'softmax')(x)
19
20 model = Model(inputs=input_padding,outputs=output)
21 model.summary()
```

Model: "functional_1"

Layer (type)	Output Shape	Param #
=====		
input_padding_ids (InputLayer)	[(None, 64, 35)]	0

lstm (LSTM)	(None, 64, 1500)	9216000

global_average_pooling1d (GlobalAveragePooling1D)	(None, 64)	0

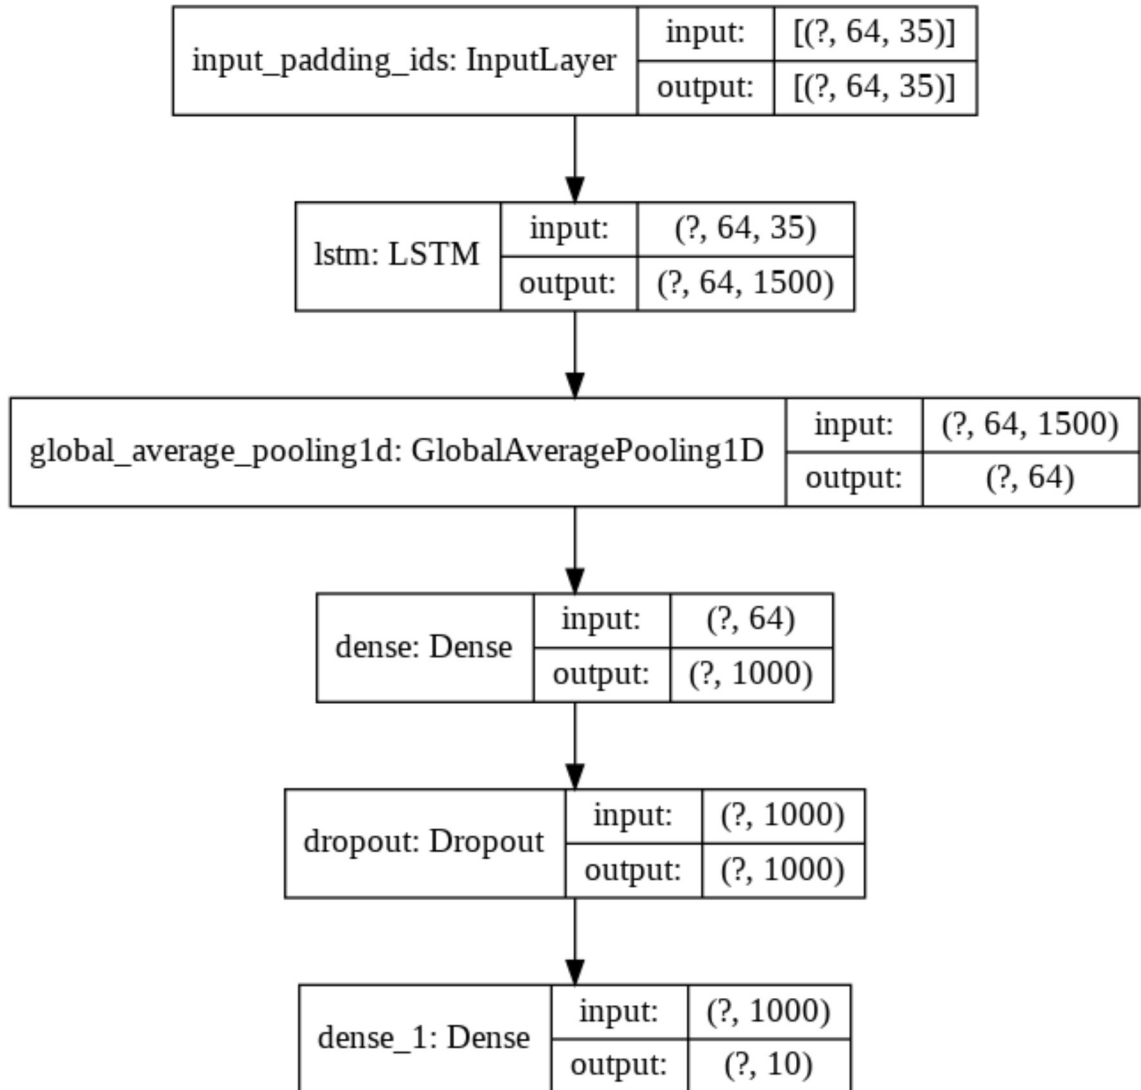
dense (Dense)	(None, 1000)	65000

dropout (Dropout)	(None, 1000)	0

dense_1 (Dense)	(None, 10)	10010
=====		
Total params: 9,291,010		
Trainable params: 9,291,010		
Non-trainable params: 0		

```
In [83]: 1 dot_img_file = '/tmp/model_1.png'
        2 tf.keras.utils.plot_model(model, to_file=dot_img_file, show_shapes=True)
```

Out[83]:



In [86]:

```
1 ACCURACY_THRESHOLD = 0.8
2 class Metrics(tf.keras.callbacks.Callback):
3     def __init__(self, validation):
4         super(Metrics, self).__init__()
5         self.validation_data = validation
6
7     def on_train_begin(self, logs={}):
8         self.val_f1s_score = []
9         self.f1_score_best = 0
10        self.epoch_value = 1
11
12
13    def on_epoch_end(self, epoch, logs={}):
14
15
16
17        val_predict = self.model.predict(self.validation_data[0])
18        val_predict = np.argmax(val_predict, axis=1)
19
20        val_targ = self.validation_data[1]
21
22        _val_f1 = f1_score(val_targ, val_predict, average='micro')
23        self.val_f1s_score.append(_val_f1)
24
25        print(' - val_f1: %f ' %(_val_f1))
26
27
28        if _val_f1 > self.f1_score_best:
29            print('F1_score improved from '+str(self.f1_score_best ) + ' to '
30                  self.f1_score_best = _val_f1
31        else:
32            print('Model not improved, still best f1_score reamins '+str(sel
33
34        self.epoch_value = self.epoch_value +1
35        if (_val_f1 >= ACCURACY_THRESHOLD):
36            print("\nReached %2.2f%% accuracy, so stopping training!!" %(ACCU
37                  self.model.stop_training = True
38            return
39
40 metrics = Metrics((X_test_spectrogram, y_test_fit))
```

```
In [87]: 1 # Call back
2
3 # Learning rate scheduler
4 def my_learning_rate(epoch, lrate):
5     if epoch <=500:
6         print('Learning rate changed to 0.0001')
7         lrate = 0.0001
8
9     else:
10        print('Learning rate changed to 0.00001')
11        lrate = 0.00001
12
13    return lrate
14
15 lrs = LearningRateScheduler(my_learning_rate)
16 earlystop = EarlyStopping(monitor='val_accuracy', min_delta=0, patience
17 filepath="/gdrive/My Drive/model_save/best_model2-{epoch:02d}.h5"
18 checkpoint = ModelCheckpoint(filepath=filepath, monitor='val_accuracy',
19 # TensorBoard Creation
20 %load_ext tensorboard
21 folder_name = datetime.datetime.now().strftime("%Y%m%d-%H%M%S")
22
23 # Create Log folder - TensorBoard
24 log_dir="/gdrive/My Drive/logs/fit/" + folder_name
25 tensorboard_callback =TensorBoard(log_dir=log_dir,histogram_freq=1, wri
```

The tensorboard extension is already loaded. To reload it, use:
%reload_ext tensorboard

```
In [88]: 1 folder_name
```

```
Out[88]: '20201122-044011'
```

```
In [ ]: 1 optim = tf.keras.optimizers.Adam(learning_rate=0.0001)
2 model.compile(loss='sparse_categorical_crossentropy', optimizer= optim,
3 model.fit(x = X_train_spectrogram, y=y_train_fit, epochs=600,verbose=1
4         callbacks =[earlystop,tensorboard_callback,metrics])
```

Epoch 1/600

1/450 [.....] - ETA: 0s - loss: 2.3066 - accuracy: 0.0625
WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/tensorflow/python/ops/summary_ops_v2.py:1277: stop (from tensorflow.python.eager.profiler) is deprecated and will be removed after 2020-07-01.

Instructions for updating:

use `tf.profiler.experimental.stop` instead.

2/450 [.....] - ETA: 1:01 - loss: 2.3039 - accuracy: 0.0625
WARNING:tensorflow:Callbacks method `on_train_batch_end` is slow compared to the batch time (batch time: 0.0130s vs `on_train_batch_end` time: 0.2597s). Check your callbacks.

450/450 [=====] - ETA: 0s - loss: 2.2577 - accuracy: 0.1379 - val_f1: 0.205000

F1_score improved from 0 to 0.205 Epoch value 1

450/450 [=====] - 16s 34ms/step - loss: 2.2577 - accuracy: 0.1379 - val_loss: 2.1771 - val_accuracy: 0.2050

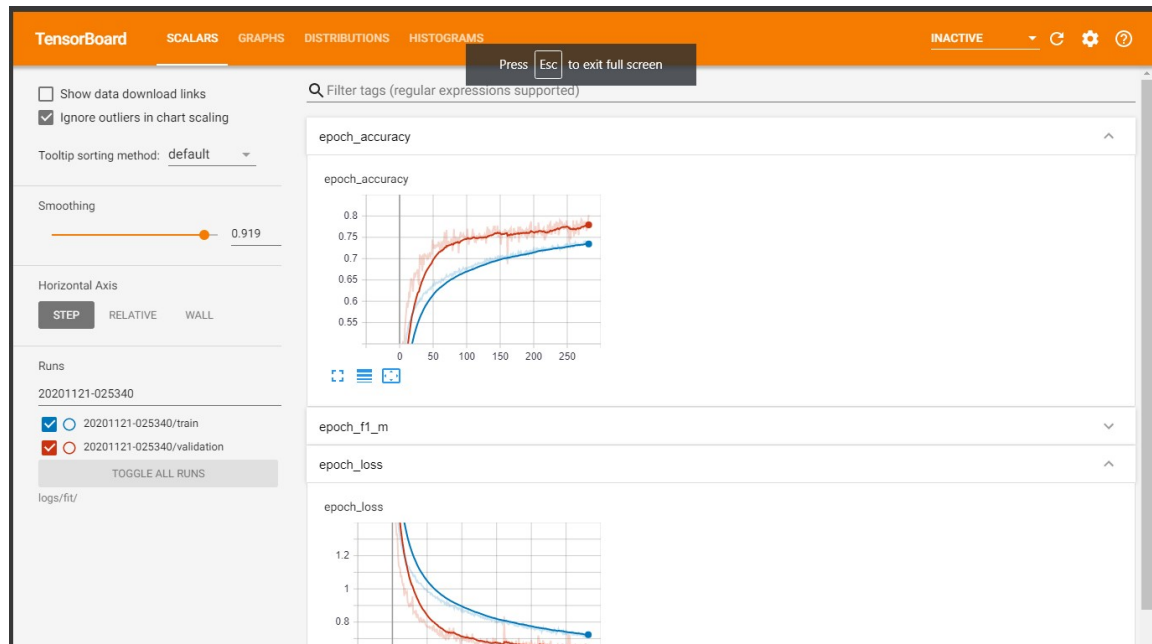
Epoch 2/600

449/450 [=====>.] - ETA: 0s - loss: 2.1640 - accuracy: 0.2062 - val_f1: 0.242500

F1_score improved from 0.205 to 0.2425 Epoch value 2

```
In [ ]: 1 #Model 1 - results
2 os.chdir('/gdrive/My Drive/')
3 %tensorboard --logdir logs/fit/
4
```

Output hidden; open in <https://colab.research.google.com> (<https://colab.research.google.com>) to view.



```
In [ ]: 1
```